

# Team9\_Report

December 16, 2021

## 1 Team 9 Project Report: Eri Kim, Arman Saduakas, Hoang Phan Pham

### 1.1 Description of Algorithm:

Our team created a classification neural network model that consists of multiple layers using different activation functions to predict a digit that is written in the image.

The activation functions used are as follows:

1. Relu: Rectified linear activation function is a linear function that outputs the input directly if it is positive, otherwise, zero. It has become the default activation function for many types of neural networks because it is easier to train and achieves better performance. One of the advantages is that its gradient is always equal to 1, which enables the model to pass the maximum amount of the error through the network during back-propagation. The relu functions are usually used as input and hidden layers because they describe variation of data since the output can be any number between 0 and positive infinity.
2. Sigmoid: Sigmoid function is one of the most widely used non-linear activation function. It transforms values between 0 and 1. It is mostly used for models when predicting probability is needed. Since probability can be only between 0 and 1, sigmoid is used in that case.
3. Softmax: Softmax function is used as an activation function in the output layer for a model that predicts a multi-class classification problems where class membership is required on more than two class labels. The network is configured to output N values, and the softmax function is used to normalize the outputs, converting each of them from weighted sum values into probabilities that sum to one. Each output of the softmax activation function can be interpreted as the probability of membership for each class label.

#### 1.1.1 Three Different Approaches to Optimize the Model

We tried to use other activation functions such as tanh; however, we decided to use sigmoid instead since our output of the model would only require 0 to 1 instead -1 to 1. We also tried adjusting the epoch size. We increased the size to optimize the model, however, increasing the epoch size started over fitting after size = 35, which resulted in decreasing the test performance. Lastly, we tried different numbers of neurons in each layer. We attempted to change the number of neurons in each layer to optimize the test performance, and the numbers we ended up with provided us the best performance, which gave us 96.63% accuracy on the test data. We also tried many different combinations of the ReLu function and various kernel initializers and found out that Random Uniform worked the best. In addition, Glorot kernel initializer performed the best with the Sigmoid activation function.

## 1.2 Imports

```
[ ]: import pandas as pd
data_train = pd.read_csv('./train.csv')
```

## 2 Data description

The dataset was obtained from Kaggle:

<https://www.kaggle.com/animatronbot/mnist-digit-recognizer>

The data has 42000 rows, each row represents 1 column for label and 784 columns for pixel of each image.

```
[ ]: print(data_train.shape)
data_train.head()
```

(42000, 785)

```
[ ]:   label  pixel0  pixel1  pixel2  pixel3  pixel4  pixel5  pixel6  pixel7  \
0      1        0        0        0        0        0        0        0        0
1      0        0        0        0        0        0        0        0        0
2      1        0        0        0        0        0        0        0        0
3      4        0        0        0        0        0        0        0        0
4      0        0        0        0        0        0        0        0        0
```

```
   pixel8  ...  pixel774  pixel775  pixel776  pixel777  pixel778  pixel779  \
0        0  ...        0        0        0        0        0        0
1        0  ...        0        0        0        0        0        0
2        0  ...        0        0        0        0        0        0
3        0  ...        0        0        0        0        0        0
4        0  ...        0        0        0        0        0        0
```

```
   pixel780  pixel781  pixel782  pixel783
0          0          0          0          0
1          0          0          0          0
2          0          0          0          0
3          0          0          0          0
4          0          0          0          0
```

[5 rows x 785 columns]

```
[ ]: import tensorflow as tf

images = data_train.drop(columns=['label'])
labels = data_train.label
labels = tf.keras.utils.to_categorical(labels)
print(images.shape)
print(labels.shape)
```

(42000, 784)

(42000, 10)

We split the data into training, validation, and testing.

Training: 60%

Validation: 15%

Test: 25%

```
[ ]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(
    images, labels, test_size=0.40, random_state=101, stratify=labels
)
x_test, x_validation, y_test, y_validation = train_test_split(
    x_test, y_test, test_size=0.375, random_state=202, stratify=y_test
)
print('Training has', y_train.shape[0], 'digits')
print('Validation has', y_validation.shape[0], 'digits')
print('Testing has', y_test.shape[0], 'digits')
```

Training has 25200 digits

Validation has 6300 digits

Testing has 10500 digits

```
[ ]: from keras.models import Sequential
from keras.layers import Dense

model = Sequential()
model.add(Dense(750, input_shape=(28 * 28,),
    ↪kernel_initializer="random_uniform", activation="relu"))
model.add(Dense(500, kernel_initializer="random_uniform", activation="relu"))
model.add(Dense(250, kernel_initializer="random_uniform", activation="relu"))
model.add(Dense(50, kernel_initializer="glorot_uniform", activation="sigmoid"))
model.add(Dense(10, kernel_initializer="he_normal", activation="softmax"))

model.compile(optimizer='sgd',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
[ ]: model_training = model.fit(
    x_train,
    y_train,
    validation_data=(x_validation, y_validation),
    epochs=35,
    batch_size=128,
)
```

Epoch 1/35

197/197 [=====] - 2s 7ms/step - loss: 0.9037 -

accuracy: 0.8043 - val\_loss: 0.5350 - val\_accuracy: 0.9081  
 Epoch 2/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.4335 -  
 accuracy: 0.9243 - val\_loss: 0.3761 - val\_accuracy: 0.9325  
 Epoch 3/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.3130 -  
 accuracy: 0.9458 - val\_loss: 0.3048 - val\_accuracy: 0.9417  
 Epoch 4/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.2456 -  
 accuracy: 0.9579 - val\_loss: 0.2510 - val\_accuracy: 0.9513  
 Epoch 5/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.2007 -  
 accuracy: 0.9664 - val\_loss: 0.2222 - val\_accuracy: 0.9548  
 Epoch 6/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.1680 -  
 accuracy: 0.9727 - val\_loss: 0.2006 - val\_accuracy: 0.9590  
 Epoch 7/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.1426 -  
 accuracy: 0.9780 - val\_loss: 0.1852 - val\_accuracy: 0.9595  
 Epoch 8/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.1215 -  
 accuracy: 0.9825 - val\_loss: 0.1785 - val\_accuracy: 0.9598  
 Epoch 9/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.1059 -  
 accuracy: 0.9852 - val\_loss: 0.1654 - val\_accuracy: 0.9627  
 Epoch 10/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0931 -  
 accuracy: 0.9881 - val\_loss: 0.1569 - val\_accuracy: 0.9654  
 Epoch 11/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0826 -  
 accuracy: 0.9901 - val\_loss: 0.1496 - val\_accuracy: 0.9654  
 Epoch 12/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0742 -  
 accuracy: 0.9911 - val\_loss: 0.1440 - val\_accuracy: 0.9660  
 Epoch 13/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0671 -  
 accuracy: 0.9923 - val\_loss: 0.1404 - val\_accuracy: 0.9648  
 Epoch 14/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.0612 -  
 accuracy: 0.9932 - val\_loss: 0.1411 - val\_accuracy: 0.9654  
 Epoch 15/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0562 -  
 accuracy: 0.9939 - val\_loss: 0.1333 - val\_accuracy: 0.9665  
 Epoch 16/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0520 -  
 accuracy: 0.9941 - val\_loss: 0.1308 - val\_accuracy: 0.9660  
 Epoch 17/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0482 -

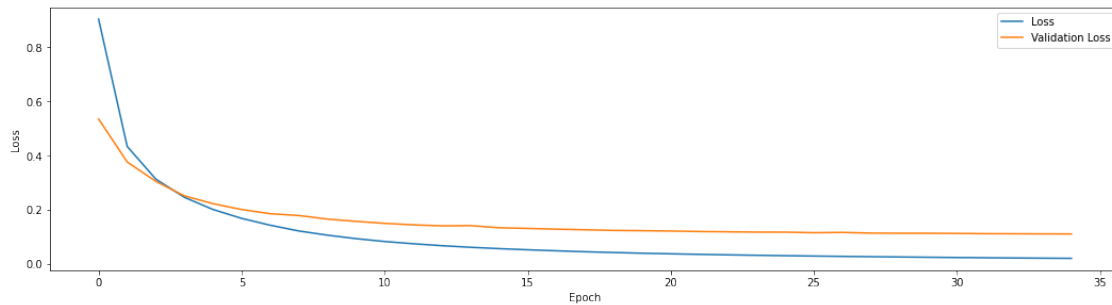
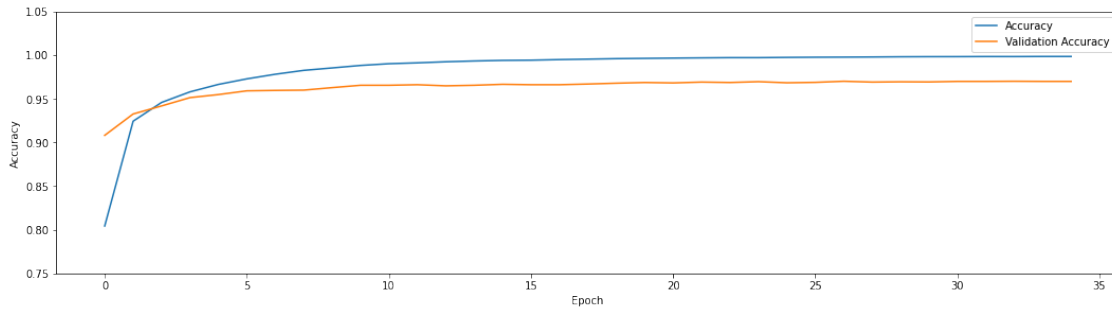
accuracy: 0.9948 - val\_loss: 0.1281 - val\_accuracy: 0.9660  
 Epoch 18/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.0447 -  
 accuracy: 0.9953 - val\_loss: 0.1260 - val\_accuracy: 0.9668  
 Epoch 19/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.0419 -  
 accuracy: 0.9959 - val\_loss: 0.1234 - val\_accuracy: 0.9678  
 Epoch 20/35  
 197/197 [=====] - 2s 8ms/step - loss: 0.0393 -  
 accuracy: 0.9962 - val\_loss: 0.1226 - val\_accuracy: 0.9686  
 Epoch 21/35  
 197/197 [=====] - 2s 9ms/step - loss: 0.0371 -  
 accuracy: 0.9965 - val\_loss: 0.1213 - val\_accuracy: 0.9681  
 Epoch 22/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.0351 -  
 accuracy: 0.9968 - val\_loss: 0.1194 - val\_accuracy: 0.9690  
 Epoch 23/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0332 -  
 accuracy: 0.9970 - val\_loss: 0.1184 - val\_accuracy: 0.9686  
 Epoch 24/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.0317 -  
 accuracy: 0.9970 - val\_loss: 0.1173 - val\_accuracy: 0.9695  
 Epoch 25/35  
 197/197 [=====] - 1s 7ms/step - loss: 0.0302 -  
 accuracy: 0.9973 - val\_loss: 0.1170 - val\_accuracy: 0.9683  
 Epoch 26/35  
 197/197 [=====] - 2s 8ms/step - loss: 0.0289 -  
 accuracy: 0.9975 - val\_loss: 0.1151 - val\_accuracy: 0.9687  
 Epoch 27/35  
 197/197 [=====] - 2s 8ms/step - loss: 0.0276 -  
 accuracy: 0.9976 - val\_loss: 0.1161 - val\_accuracy: 0.9700  
 Epoch 28/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0265 -  
 accuracy: 0.9978 - val\_loss: 0.1136 - val\_accuracy: 0.9690  
 Epoch 29/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0254 -  
 accuracy: 0.9980 - val\_loss: 0.1131 - val\_accuracy: 0.9694  
 Epoch 30/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0244 -  
 accuracy: 0.9981 - val\_loss: 0.1130 - val\_accuracy: 0.9692  
 Epoch 31/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0234 -  
 accuracy: 0.9982 - val\_loss: 0.1124 - val\_accuracy: 0.9698  
 Epoch 32/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0225 -  
 accuracy: 0.9983 - val\_loss: 0.1115 - val\_accuracy: 0.9698  
 Epoch 33/35  
 197/197 [=====] - 1s 6ms/step - loss: 0.0217 -

```
accuracy: 0.9983 - val_loss: 0.1111 - val_accuracy: 0.9700
Epoch 34/35
197/197 [=====] - 1s 6ms/step - loss: 0.0210 -
accuracy: 0.9984 - val_loss: 0.1108 - val_accuracy: 0.9698
Epoch 35/35
197/197 [=====] - 1s 6ms/step - loss: 0.0204 -
accuracy: 0.9984 - val_loss: 0.1105 - val_accuracy: 0.9698
```

```
[ ]: import matplotlib.pyplot as plt

_, axis = plt.subplots(figsize=(18, 10))
axis = plt.subplot(2, 1, 1)
axis.plot(
    range(len(model_training.history.get("accuracy"))),
    model_training.history.get("accuracy"),
    label="Accuracy",
)
axis.plot(
    range(len(model_training.history.get("val_accuracy"))),
    model_training.history.get("val_accuracy"),
    label="Validation Accuracy",
)
axis = plt.gca()
axis.set_ylim([0.75, 1.05])
axis.legend(loc="best")
axis.set_xlabel("Epoch")
axis.set_ylabel("Accuracy")
plt.show()

_, axis = plt.subplots(figsize=(18, 10))
axis = plt.subplot(2, 1, 2)
axis.plot(
    range(len(model_training.history.get("loss"))),
    model_training.history.get("loss"),
    label="Loss"
)
axis.plot(
    range(len(model_training.history.get("val_loss"))),
    model_training.history.get("val_loss"),
    label="Validation Loss",
)
axis.legend(loc="best")
axis.set_xlabel("Epoch")
axis.set_ylabel("Loss")
plt.show()
```



```
[ ]: prediction = model.predict(x_test)
scores = model.evaluate(x_test, y_test, verbose=0)
accuracy = scores[1] * 100
error = 100 - scores[1] * 100
print("Accuracy: %.2f%%" % accuracy)
print("Error: %.2f%%" % error)
```

Accuracy: 96.63%  
Error: 3.37%

```
[ ]: import numpy as np

predicted, actual = [], []
for p in prediction:
    max = float(0)
    index, max_i = 0, -1
    for n in p:
        if float(n) > max:
            max = float(n)
            max_i = index
        index += 1
    predicted.append(max_i)

for digit in y_test:
```

```

actual.append(np.argmax(digit))

y_actual = pd.Series(actual, name="Actual")
y_predict = pd.Series(predicted, name="Predicted")

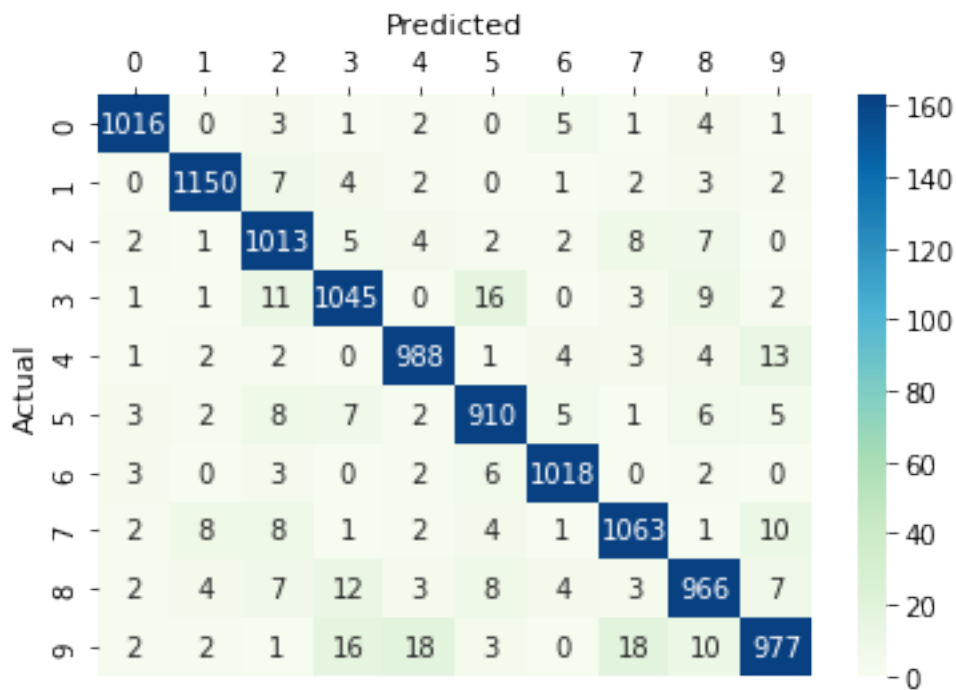
```

```

[ ]: import seaborn as sns
from sklearn.metrics import confusion_matrix

cf_matrix = confusion_matrix(y_actual, y_predict)
axis = sns.heatmap(cf_matrix, vmin=0, vmax=163, annot=True, fmt="d",
    ↳linewidths=0, cmap="GnBu")
axis.xaxis.tick_top()
axis.set_title('Predicted', size=11)
axis.set_ylabel('Actual', loc="center", size=11)
plt.show()

```



```

[ ]: count = 0
for i in range(len(prediction)):
    if np.argmax(prediction[i]) != np.argmax(y_test[i]):
        count += 1
print("Number of wrong classified digits: " + str(count))

```

Number of wrong classified digits: 354